

Problem set 1: Probability

Due October 2, 2023, at 9pm

(Your name here)

NOTE: Start with the file `ps1_2023_probability.qmd` (available from the github repository at <https://github.com/UChicago-pol-methods/IntroQSS-F23/tree/main/assignments>). Modify that file to include your answers. Make sure you can “knit” the file (e.g. in RStudio by clicking on the *Knit* button). Submit both the `qmd` file and the knitted PDF via Canvas.

Problem 1: joint probability, conditional probability, independence

Suppose a speech is chosen at random from a collection of politicians’ speeches. Let A denote the event that the speech contains the word “congratulate”. Let B denote the event that the speech contains the word “shameful”. Let $P(A)$ denote the (marginal) probability of event A and $P(B)$ denote the (marginal) probability of event B .

(1a) Suppose events A and B are independent. Fill out the joint probability table below using $P(A)$ and $P(B)$.

Result	Probability
$A \cap B$	
$A^C \cap B$	
$A \cap B^C$	
$A^C \cap B^C$	

(1b) Do you think events A and B really would be independent, given a collection of speeches from e.g. the US Congress or the British House of Commons? Say why or why not, perhaps with reference to the definition of independence.

Answer:

Now suppose the true joint probabilities are:

Result	Probability
$A \cap B$	1/8
$A^C \cap B$	1/4
$A \cap B^C$	7/24
$A^C \cap B^C$	1/3

(1c) What are $P(A)$ and $P(B)$, i.e. the marginal probabilities of A and B ?

Answer:

(1d) What is $P(A|B)$, i.e. the conditional probability of A given B ? What would it be if the events were independent?

Answer:

(1e) How much more or less likely is a speech to contain both “congratulate” and “shameful” than would be the case if the events were independent?

Answer:

Problem 2: R coding

First set your seed to 123 so that our answers are comparable.

```
set.seed(123)
```

(2a) Create a vector of length 1000 that resembles a sample of speeches in Problem 1, where 1 indicates that the word “congratulate” appears in the speech and 0 indicates that it does not. Use the marginal probability from problem (1c). Store this vector in a variable called `A_vec`. Report `mean(A)`.

Answer:

```
# your code goes here.
```

(2b) Do the same for the event that the word “shameful” appears in the speech. That is, create a vector of length 1000, where 1 indicates that the word “shameful” appears in the speech and 0 indicates that it does not. Use the marginal probability from problem (1c). Store this vector in a variable called `B_vec`. Report `mean(B_vec)`. (Do not reset your seed.)

Answer:

Suppose that your vectors `A` and `B` are measurements relating to the same 1000 speeches – that is, if the first element of both is 1, that means both “shameful” and “congratulate” appeared in the first speech.

(2c) Restricting to speeches in which “shameful” appears, in what proportion did “congratulate” appear?

Answer:

(2d) Answer (2a)-(2c) again, but this time for a sample of 1 million speeches. Compare your answers with the smaller and larger sample.

Answer:

(2e) If our sample was infinitely large, what would the answer to (2c) be? How does this relate to independence and conditional probability?

Answer: